

# Working With Imbalanced Data

## Python Scikit-Learn

Ben Brock

# Increasing successful detections using data resampling

- One approach to handle class imbalance
  - Undersampling
  - Oversampling
  - SMOTE
- Another approach to handle class imbalance
  - Scikit-learn models class weight option
- Find the optimal machine learning model
  - GridSearchCV

# Imbalance-Learn

<http://imbalanced-learn.org>

```
pip install -U imbalanced-learn
```

Extends sklearn API

# Example Sampler

To resample a data sets, each sampler implements:

```
data_resampled, targets_resampled = obj.sample(data, targets)
```

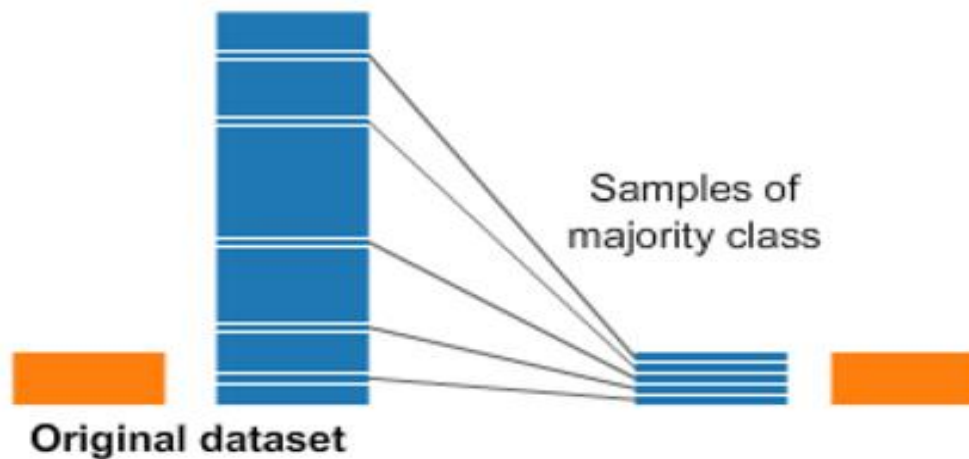
Fitting and sampling can also be done in one step:

```
data_resampled, targets_resampled = obj.fit_sample(data, targets)
```

In Pipelines: Sampling only done in `fit`!

# Undersampling

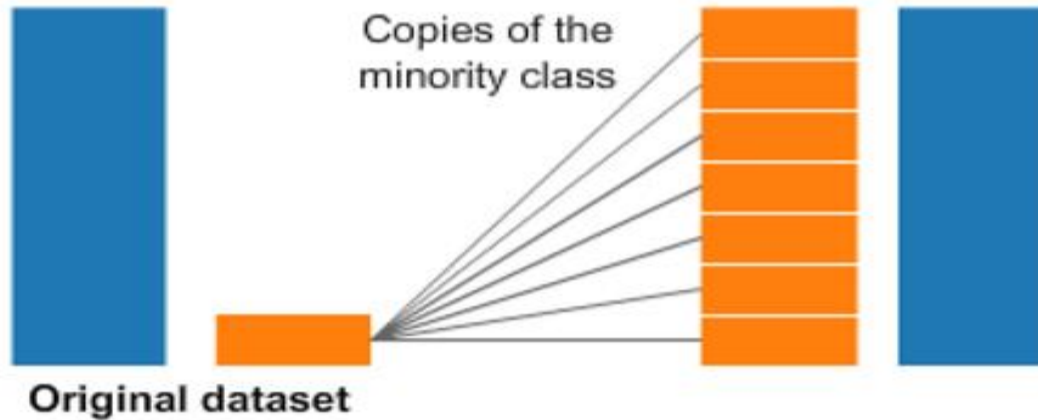
- Undersampling



- Drawback: Throwing away a lot of good data and information

# Oversampling

- Oversampling

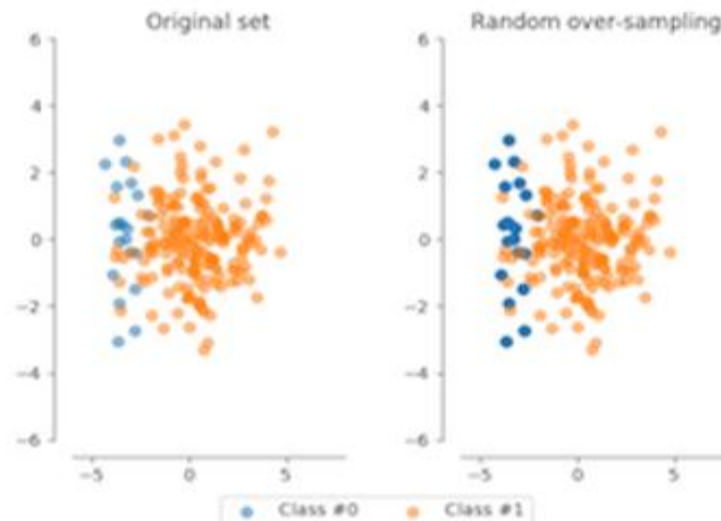


- Drawback: Training on duplicate data

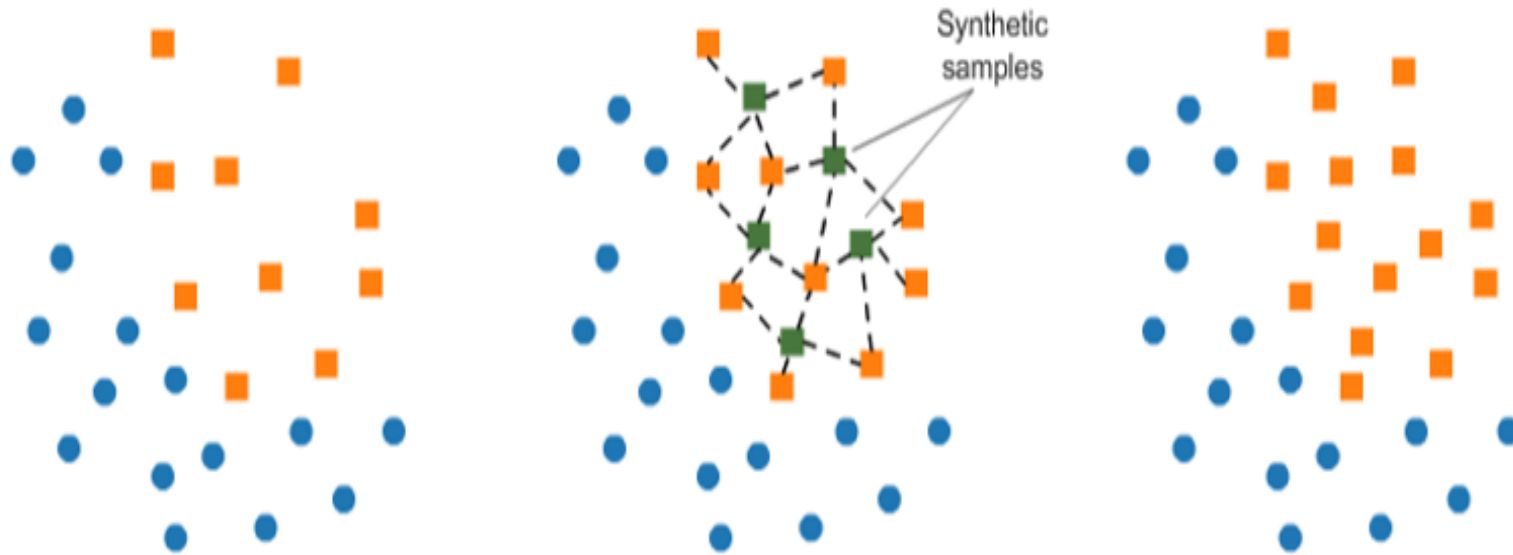
# Oversampling in Python

- Compatible with scikit-learn

```
from imblearn.over_sampling import RandomOverSampler  
  
method = RandomOverSampler()  
X_resampled, y_resampled = method.fit_sample(X, y)  
  
compare_plots(X_resampled, y_resampled, X, y)
```



# Synthetic Minority Oversampling Technique (SMOTE)



Source: <https://www.kaggle.com/rafjaa/resampling-strategies-for-imbalanced-datasets>



# Synthetic Minority Oversampling Technique (SMOTE)

- Oversampling the minority observations
  - Not just copying the minority classes, instead SMOTE uses characteristics of KNN nearest neighbors of old cases to create new synthetic cases, and thereby avoids creating duplications

# Which resampling method to use?

- Random Under Sampling (RUS): throw away data, computationally efficient
  - If you have LARGE amount of data and many categorical type (i.e., fraud) cases
- Random Over Sampling (ROS): straightforward and simple, but training your model on many duplicates
- Synthetic Minority Oversampling Technique (SMOTE): more sophisticated and realistic dataset, but you are training on "fake" data

# When to use resampling methods

- Use resampling methods on your training set, never on your test set!

```
# Define resampling method and split into train and test
method = SMOTE(kind='borderline1')
X_train, X_test, y_train, y_test = train_test_split(X, y,
    train_size=0.8, random_state=0)

# Apply resampling to the training data only
X_resampled, y_resampled = method.fit_sample(X_train, y_train)

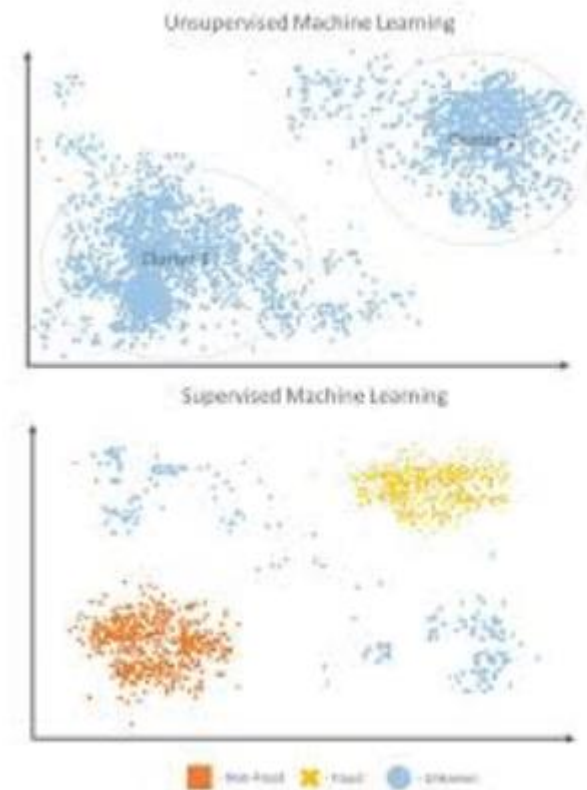
# Continue fitting the model and obtain predictions
model = LogisticRegression()
model.fit(X_resampled, y_resampled)

# Get your performance metrics
predicted = model.predict(X_test)
print(classification_report(y_test, predicted))
```

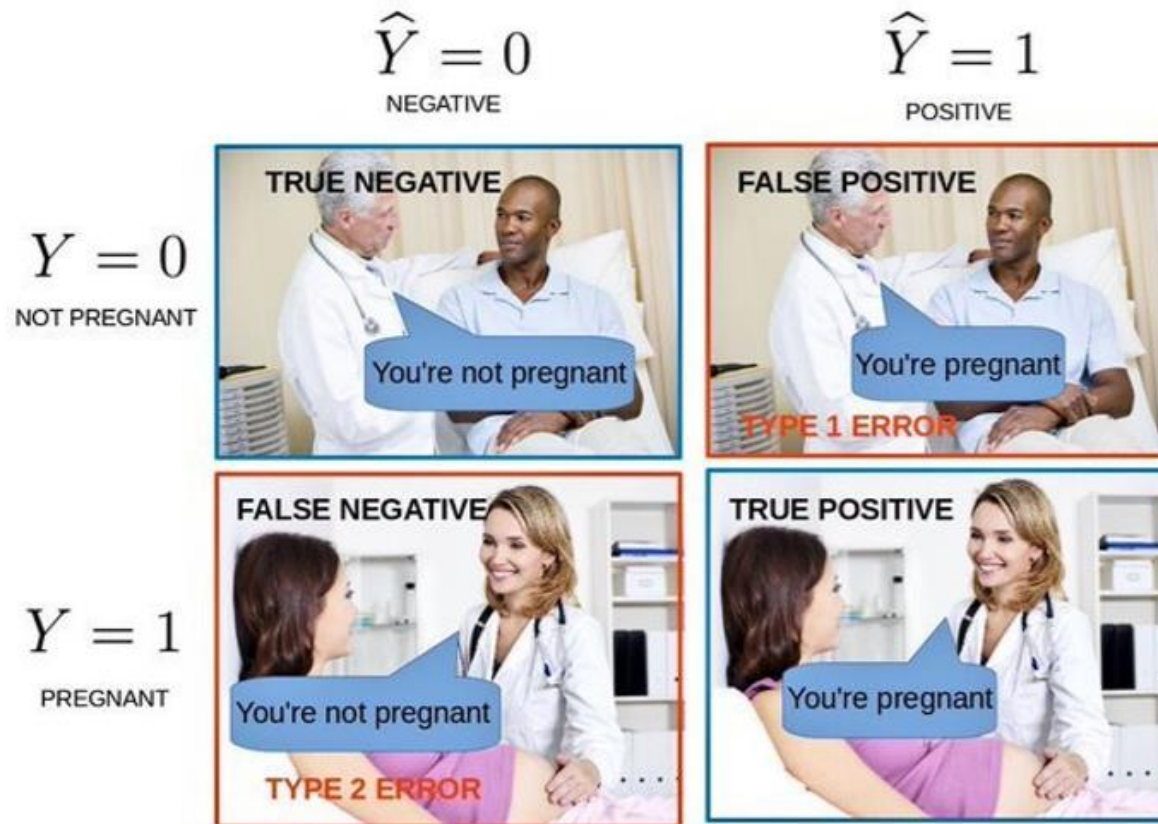
- Goal is to better train your model and give it more balanced data!!!

# Why use machine learning

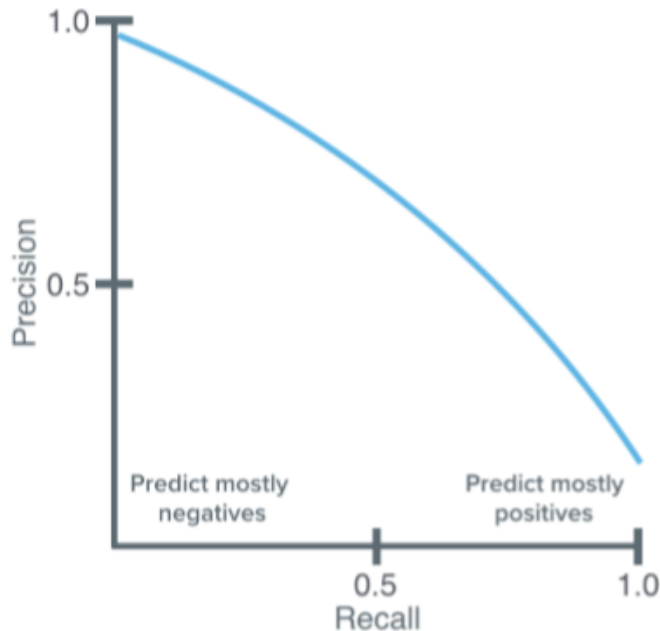
- Machine learning models adapt to the data, and thus can change over time
- Uses all the data combined rather than a threshold per feature
- Can give a score, rather than a yes/no
- Will typically have a better performance and can be combined with rule



# False positives, False negatives



# Precision Recall Trade-Off



$$Precision = \frac{\#True\ Positives}{\#True\ Positives + \#False\ Positives}$$

$$Recall = \frac{\#True\ Positives}{\#True\ Positives + \#False\ Negatives}$$

$$\begin{aligned} F - measure &= \frac{2 \times Precision \times Recall}{Precision + Recall} \\ &= \frac{2 \times TP}{2 \times TP + FP + FN} \end{aligned}$$

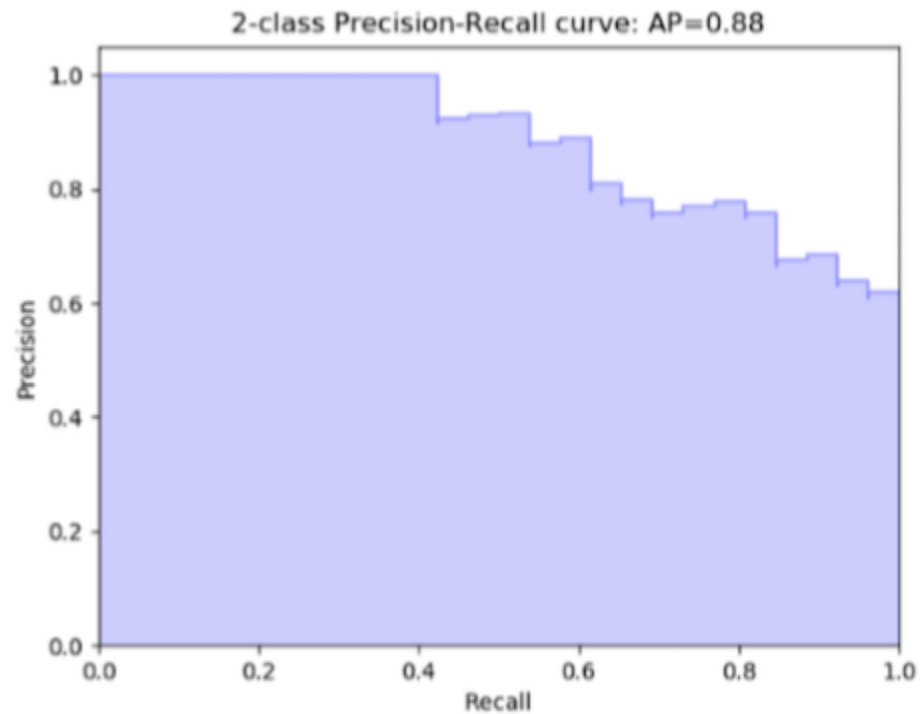
# Obtaining performance metrics

```
# Import the packages
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import average_precision_score

# Calculate average precision and the PR curve
average_precision = average_precision_score(y_test, predicted)

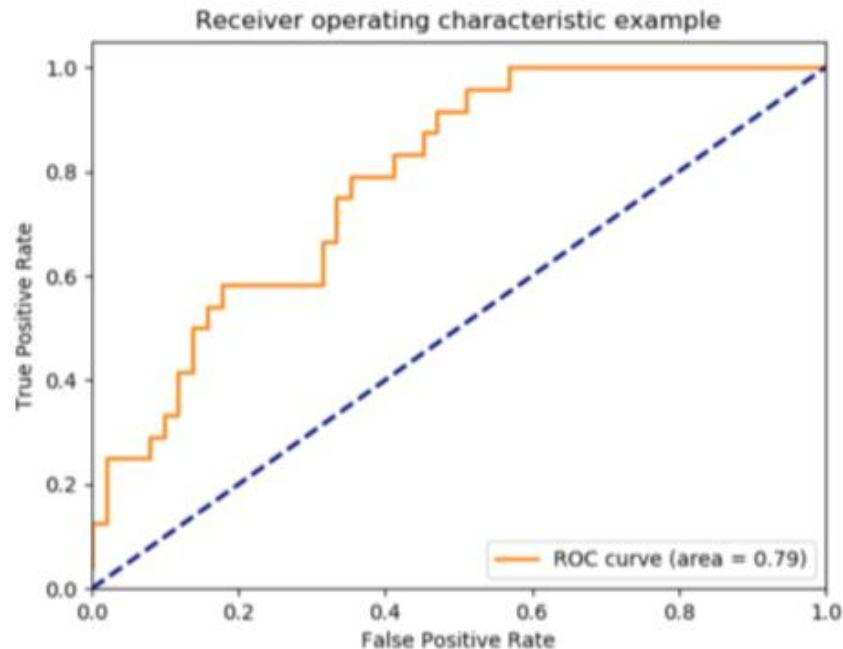
# Obtain precision and recall
precision, recall, _ = precision_recall_curve(y_test, predicted)
```

# Precision-Recall Curve





# ROC curve to compare algorithms



```
# Obtain model probabilities
probs = model.predict_proba(X_test)

# Print ROC_AUC score using probabilities
print(metrics.roc_auc_score(y_test, probs[:, 1]))
```

# Confusion Matrix and Classification Report

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
# Obtain predictions
predicted = model.predict(X_test)
```

```
# Print classification report using predictions
print(classification_report(y_test, predicted))
```

	precision	recall	f1-score	support	
	0.0	0.99	1.00	1.00	2099
	1.0	0.96	0.80	0.87	91
avg / total		0.99	0.99	0.99	2190

```
# Print confusion matrix using predictions
print(confusion_matrix(y_test, predicted))
```

```
[[2096   3]
 [  18  73]]
```

# Balance Weights

```
model = RandomForestClassifier(class_weight='balanced')
```

```
model = RandomForestClassifier(class_weight='balanced_subsample')
```

```
model = LogisticRegression(class_weight='balanced')
```

```
model = SVC(kernel='linear', class_weight='balanced', probability=True)
```

# Hyperparameter tuning

```
model = RandomForestClassifier(class_weight={0:1,1:4},random_state=1)
```

```
model = LogisticRegression(class_weight={0:1,1:4}, random_state=1)
```

```
model = RandomForestClassifier(n_estimators=10,  
    criterion='gini',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    max_features='auto',  
    n_jobs=-1, class_weight=None)
```

# Using GridSearchCV

```
from sklearn.model_selection import GridSearchCV
```

```
# Create the parameter grid
param_grid = {
    'max_depth': [80, 90, 100, 110],
    'max_features': [2, 3],
    'min_samples_leaf': [3, 4, 5],
    'min_samples_split': [8, 10, 12],
    'n_estimators': [100, 200, 300, 1000]
}
```

```
# Define which model to use
model = RandomForestRegressor()
```

```
# Instantiate the grid search model
grid_search_model = GridSearchCV(estimator = model,
    param_grid = param_grid, cv = 5,
    n_jobs = -1, scoring='f1')
```

# Finding the best model with GridSearchCV

```
# Fit the grid search to the data  
grid_search_model.fit(X_train, y_train)
```

```
# Get the optimal parameters  
grid_search_model.best_params_
```

```
{'bootstrap': True,  
 'max_depth': 80,  
 'max_features': 3,  
 'min_samples_leaf': 5,  
 'min_samples_split': 12,  
 'n_estimators': 100}
```

```
# Get the best_estimator results  
grid_search.best_estimator_  
grid_search.best_score_
```

# Working with Imbalanced Data

- Worked with highly imbalanced data
- Learned how to resample your data
- Learned about different resampling methods

# Detailed



# Changing Thresholds

```
data = load_breast_cancer()

X_train, X_test, y_train, y_test = train_test_split(
    data.data, data.target, stratify=data.target, random_state=0)

lr = LogisticRegression().fit(X_train, y_train)
y_pred = lr.predict(X_test)

classification_report(y_test, y_pred)
```

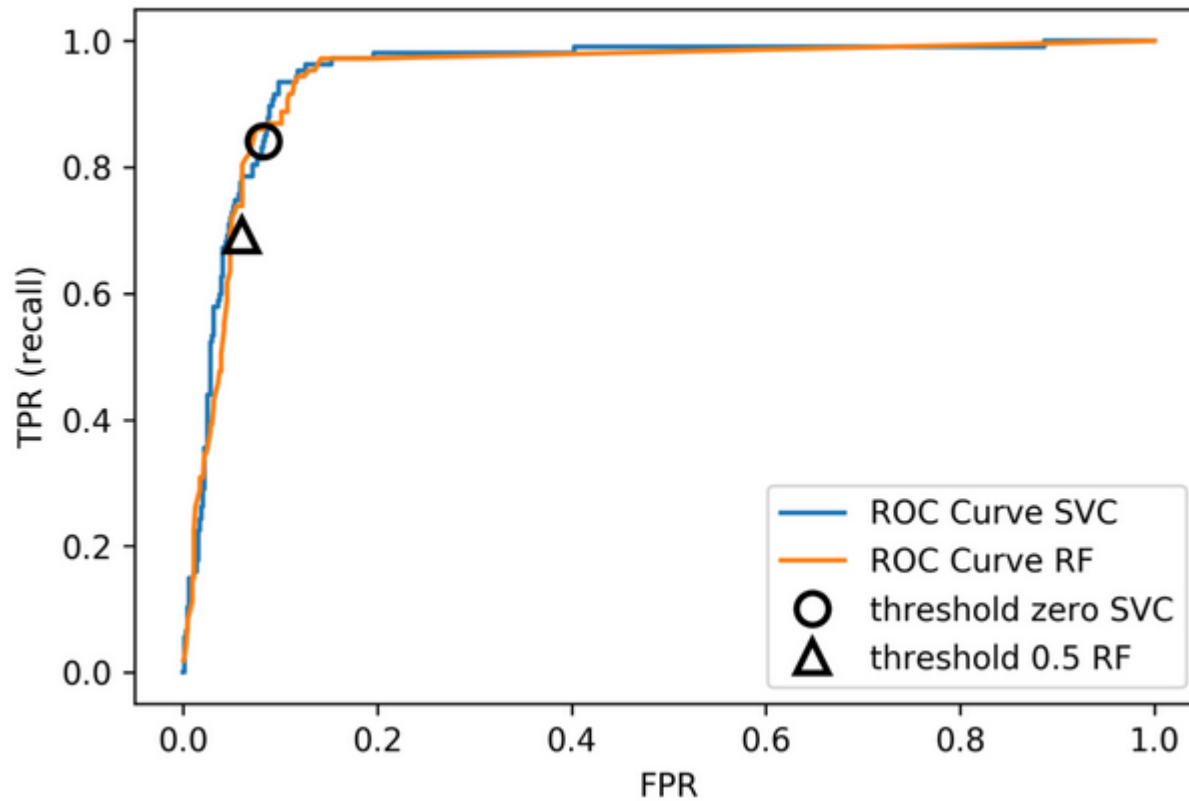
	precision	recall	f1-score	support
0	0.91	0.92	0.92	53
1	0.96	0.94	0.95	90
avg/total	0.94	0.94	0.94	143

```
y_pred = lr.predict_proba(X_test)[:, 1] > .85

classification_report(y_test, y_pred)
```

	precision	recall	f1-score	support
0	0.84	1.00	0.91	53
1	1.00	0.89	0.94	90
avg/total	0.94	0.93	0.93	143

# ROC Curve



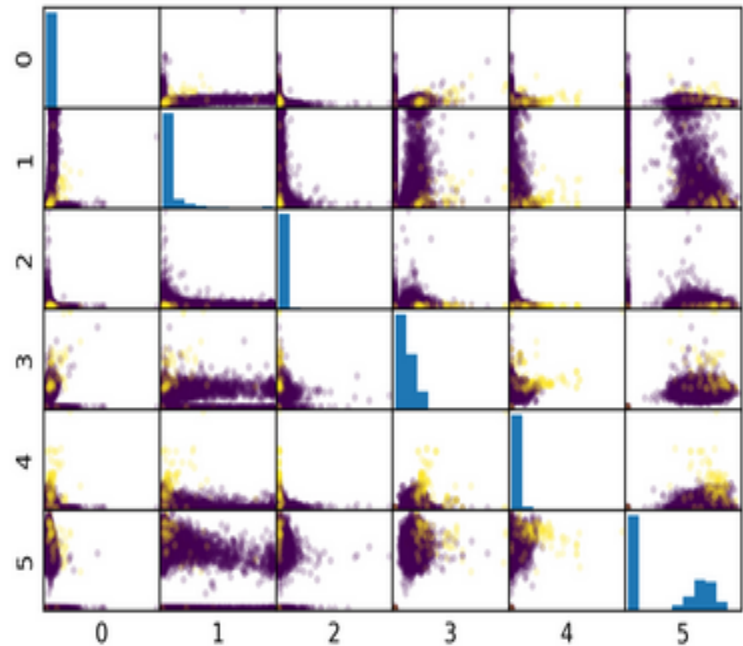
# Mammography Data

```
from sklearn.datasets import fetch_openml
# mammography https://www.openml.org/d/310
data = fetch_openml('mammography')
X, y = data.data, data.target
y = (y.astype(np.int) + 1) // 2
X.shape
```

(11183, 6)

```
np.bincount(y)
```

array([10923, 260])



# Mammography Data

```
from sklearn.model_selection import cross_validate
from sklearn.linear_model import LogisticRegression

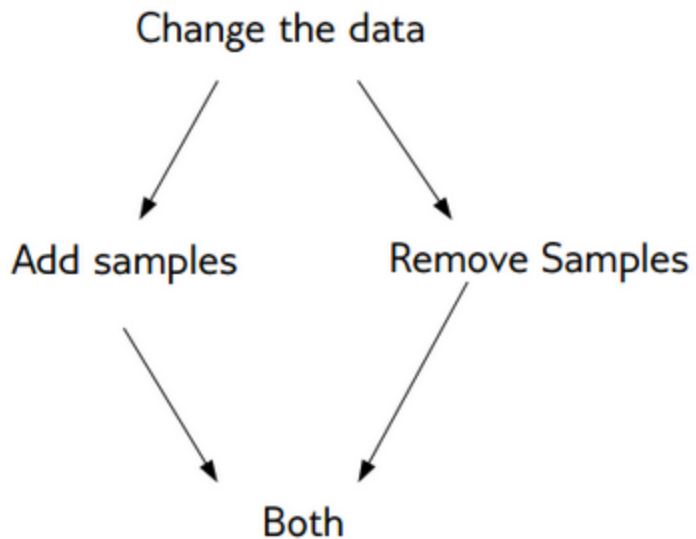
scores = cross_validate(LogisticRegression(),
                        X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()
```

0.920, 0.630

```
from sklearn.ensemble import RandomForestClassifier
scores = cross_validate(RandomForestClassifier(n_estimators=100),
                        X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()
```

0.939, 0.722

# Basic Approaches



Change the training procedure

# Sampler

To resample a data sets, each sampler implements:

```
data_resampled, targets_resampled = obj.sample(data, targets)
```

Fitting and sampling can also be done in one step:

```
data_resampled, targets_resampled = obj.fit_sample(data, targets)
```

In Pipelines: Sampling only done in `fit`!

# Random Undersampling

```
from imblearn.under_sampling import RandomUnderSampler
rus = RandomUnderSampler(replacement=False)
X_train_subsample, y_train_subsample = rus.fit_sample(
    X_train, y_train)
print(X_train.shape)
print(X_train_subsample.shape)
print(np.bincount(y_train_subsample))
```

(8387, 6)

(390, 6)

[195 195]

# Random Undersampling

```
from imblearn.pipeline import make_pipeline as make_imb_pipeline

undersample_pipe = make_imb_pipeline(RandomUnderSampler(), LogisticRegressionCV())
scores = cross_validate(undersample_pipe,
                        X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()
# baseline was 0.920, 0.630
```

0.927, 0.527

```
undersample_pipe_rf = make_imb_pipeline(RandomUnderSampler(),
                                         RandomForestClassifier())
scores = cross_validate(undersample_pipe_rf,
                        X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()
# baseline was 0.939, 0.722
```



# Random Oversampling

```
from imblearn.over_sampling import RandomOverSampler
ros = RandomOverSampler()
X_train_oversample, y_train_oversample = ros.fit_sample(
    X_train, y_train)
print(X_train.shape)
print(X_train_oversample.shape)
print(np.bincount(y_train_oversample))
```

(8387, 6)

(16384, 6)

[8192 8192]

# Random Oversampling

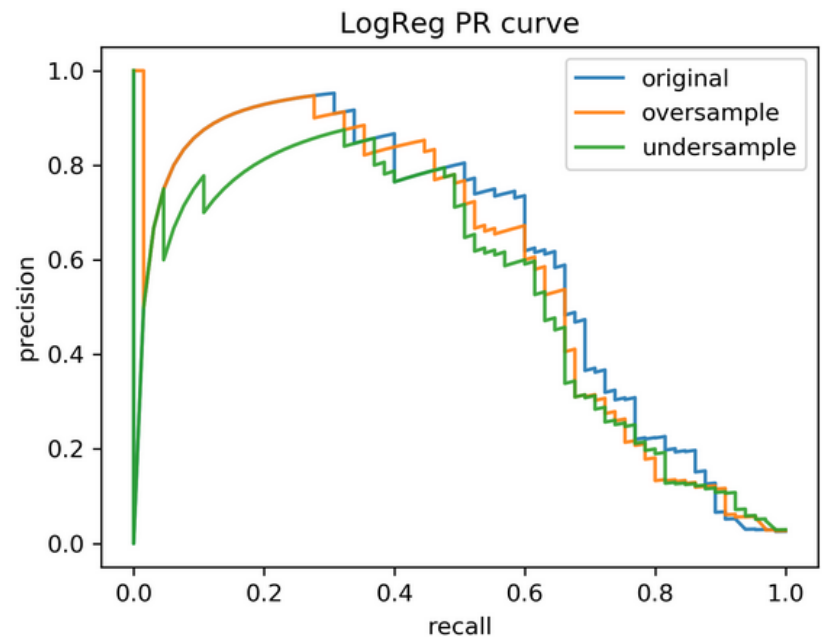
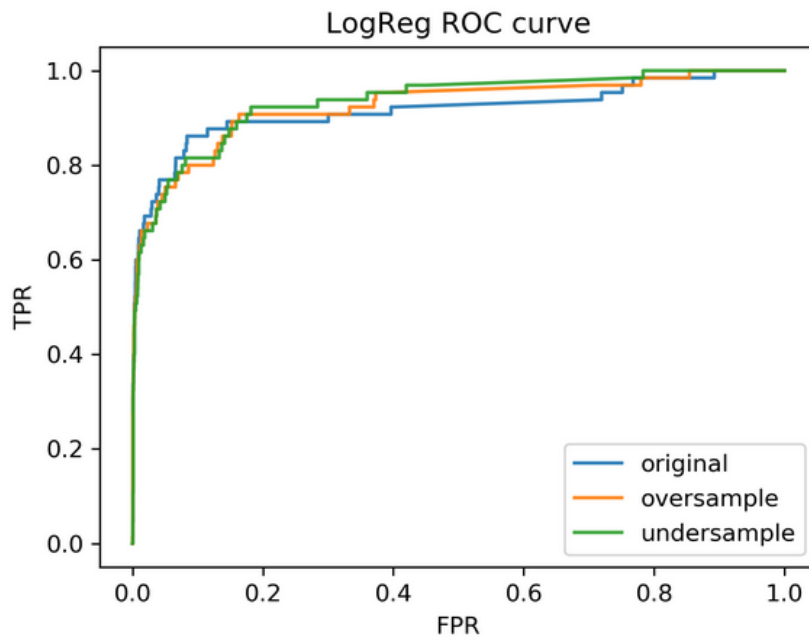
```
oversample_pipe = make_imb_pipeline(RandomOverSampler(), LogisticRegression())
scores = cross_validate(oversample_pipe,
                        X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()
# baseline was 0.920, 0.630
```

0.917, 0.585

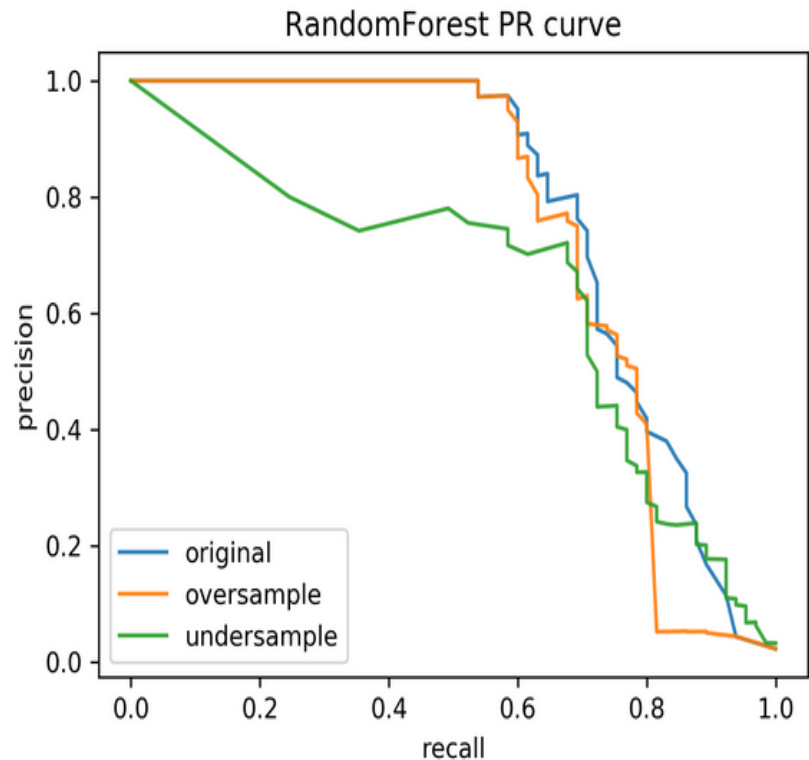
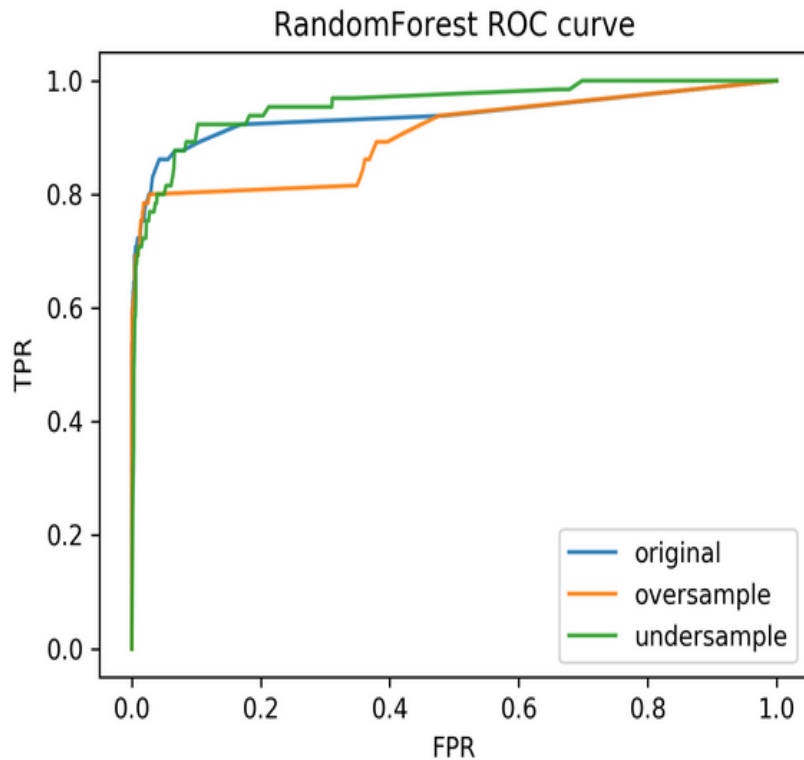
```
oversample_pipe_rf = make_imb_pipeline(RandomOverSampler(),
                                       RandomForestClassifier())
scores = cross_validate(oversample_pipe_rf,
                        X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()
# baseline was 0.939, 0.722
```

0.926, 0.715

# Curves for LogReg



# Curves for Random Forest



# ROC or PR?

FPR or Precision?

$$\text{FPR} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

# Class-weights

- Instead of repeating samples, re-weight the loss function
- Works for most models!
- Same effect as over-sampling (though not random), but not as expensive (dataset size the same).

# Class-weights in linear models

$$\min_{w \in \mathbb{R}^p, b \in \mathbb{R}} -C \sum_{i=1}^n \log(\exp(-y_i(w^T \mathbf{x}_i + b)) + 1) + \|w\|_2^2$$

$$\min_{w \in \mathbb{R}^p, b \in \mathbb{R}} -C \sum_{i=1}^n c_{y_i} \log(\exp(-y_i(w^T \mathbf{x}_i + b)) + 1) + \|w\|_2^2$$

Similar for linear and non-linear SVM

# Class weights in trees

Gini Index:

$$H_{\text{gini}}(X_m) = \sum_{k \in \mathcal{Y}} p_{mk}(1 - p_{mk})$$

$$H_{\text{gini}}(X_m) = \sum_{k \in \mathcal{Y}} c_k p_{mk}(1 - p_{mk})$$

Prediction:

Weighted vote



# Using Class-Weights

```
scores = cross_validate(LogisticRegression(class_weight='balanced'),  
                        X_train, y_train, cv=10,  
                        scoring=('roc_auc', 'average_precision'))  
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()  
# baseline was 0.920, 0.630
```

0.918, 0.587

```
scores = cross_validate(RandomForestClassifier(n_estimators=100,  
                                              class_weight='balanced'),  
                        X_train, y_train, cv=10,  
                        scoring=('roc_auc', 'average_precision'))  
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()  
# baseline was 0.939, 0.722
```

0.917, 0.701

# Ensemble Resampling

- Random resampling separate for each instance in an ensemble!
- Chen, Liaw, Breiman: “Using random forest to learn imbalanced data.”
- Paper: “Exploratory Undersampling for Class Imbalance Learning”
- Not in sklearn (yet)
- Easy with imblearn

# Easy Ensemble with imblearn

```
from sklearn.tree import DecisionTreeClassifier
from imblearn.ensemble import BalancedBaggingClassifier

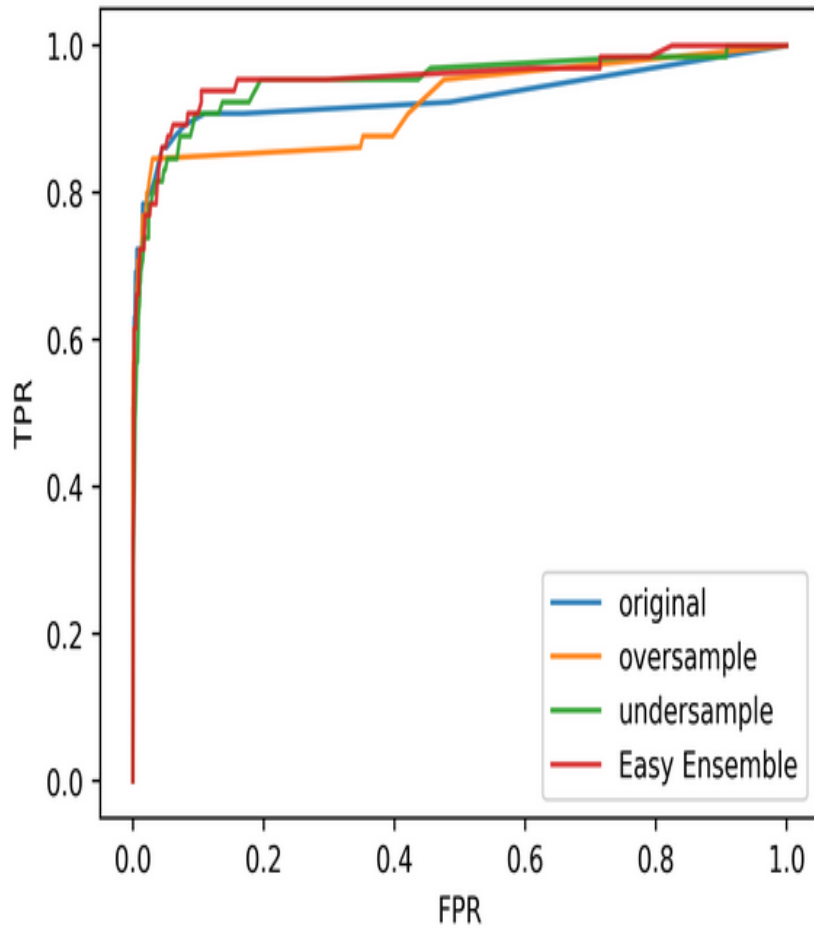
# from imblearn.ensemble import BalancedRandomForestClassifier
# resampled_rf = BalancedRandomForestClassifier()

tree = DecisionTreeClassifier(max_features='auto')
resampled_rf = BalancedBaggingClassifier(base_estimator=tree,
                                         n_estimators=100, random_state=0)

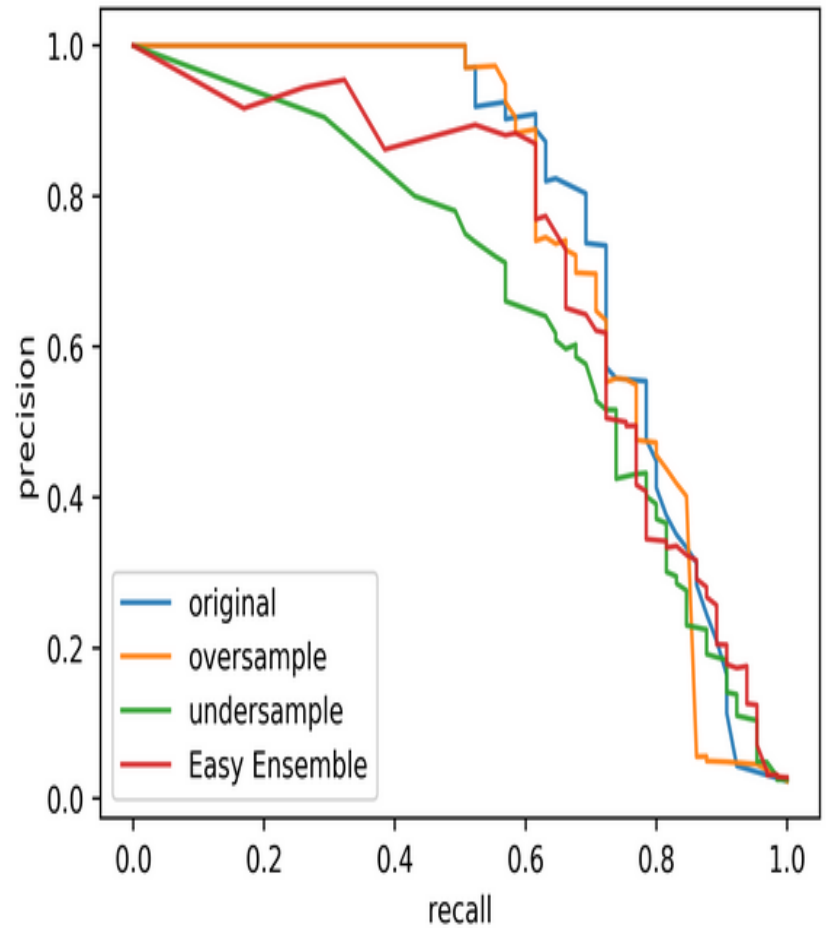
scores = cross_validate(resampled_rf,
                        X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
scores['test_roc_auc'].mean(), scores['test_average_precision'].mean()
# baseline was 0.939, 0.722
```

0.957, 0.654

RandomForest ROC curve



RandomForest PR curve

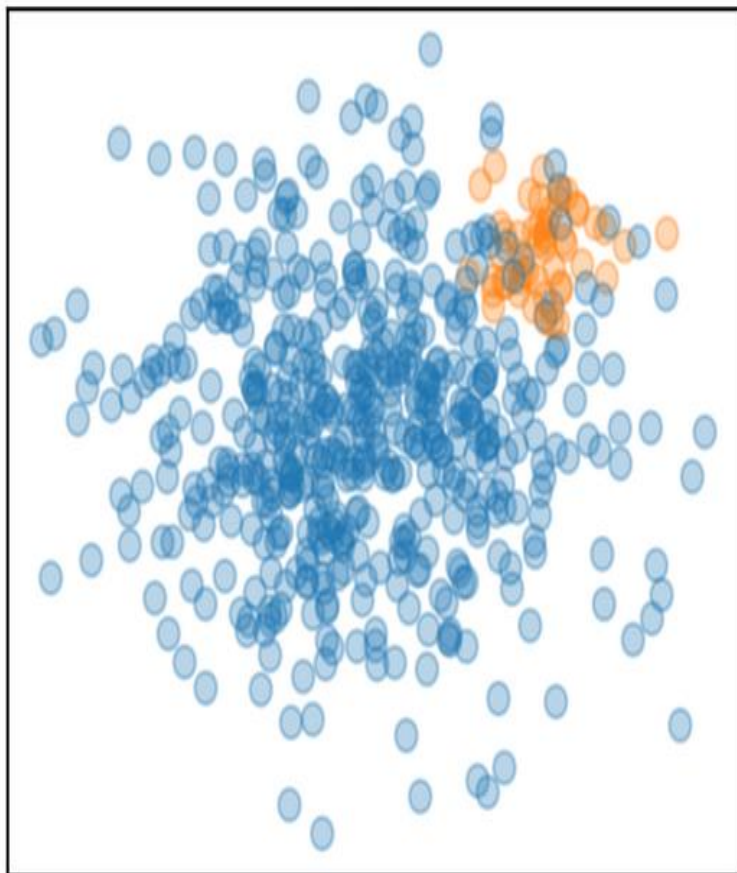


# **Synthetic Sample Generation**

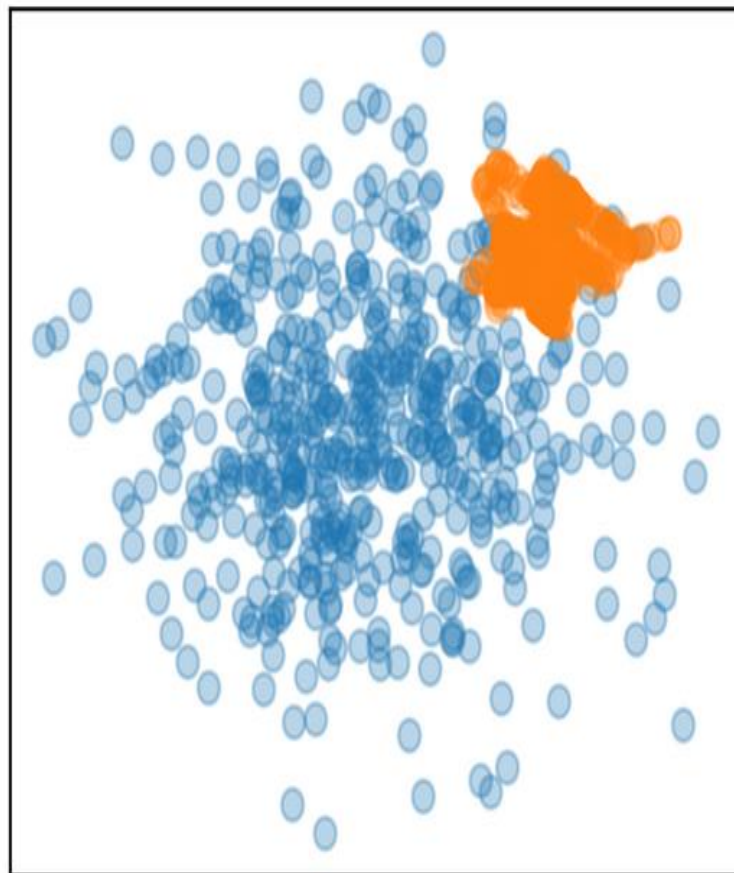
# Synthetic Minority Oversampling Technique (SMOTE)

- Adds synthetic interpolated data to smaller class
- For each sample in minority class:
  - Pick random neighbor from  $k$  neighbors.
  - Pick point on line connecting the two uniformly (or within rectangle)
  - Repeat

Original



SMOTE



# SMOTE ...

```
smote_pipe = make_imb_pipeline(SMOTE(), LogisticRegression())
scores = cross_validate(smote_pipe, X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
pd.DataFrame(scores)[['test_roc_auc', 'test_average_precision']].mean()
# baseline was 0.920, 0.630
```

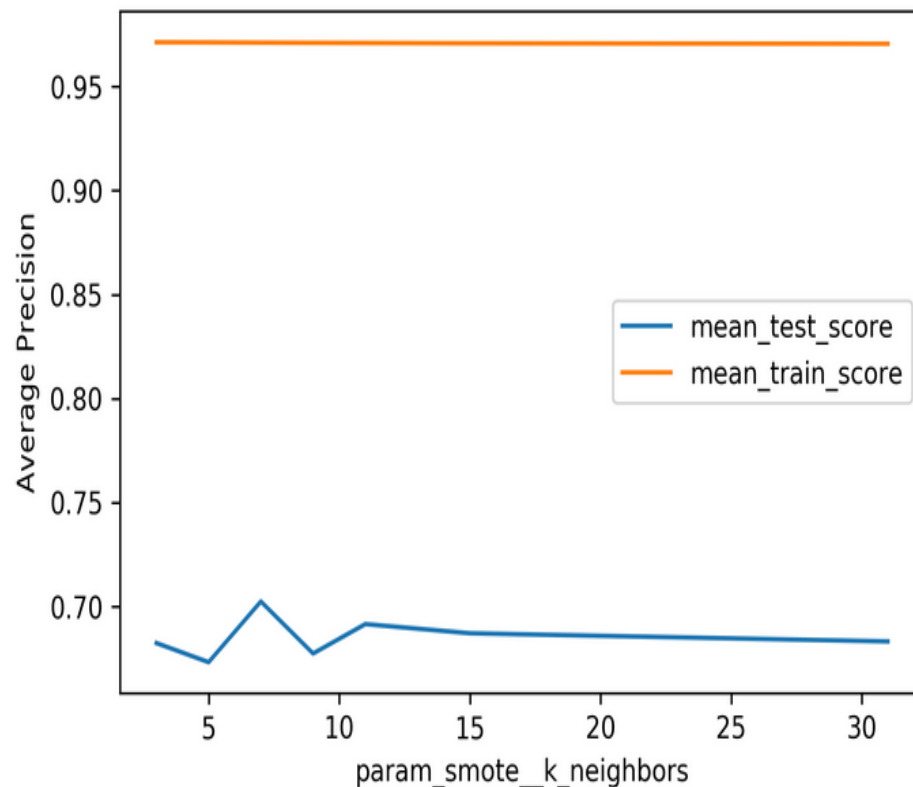
0.919, 0.585

```
smote_pipe_rf = make_imb_pipeline(SMOTE(),
                                   RandomForestClassifier(n_estimators=100))
scores = cross_validate(smote_pipe_rf, X_train, y_train, cv=10,
                        scoring=('roc_auc', 'average_precision'))
pd.DataFrame(scores)[['test_roc_auc', 'test_average_precision']].mean()
# baseline was 0.939, 0.722
```

0.946, 0.688



```
param_grid = {'smote__k_neighbors': [3, 5, 7, 9, 11, 15, 31]}
search = GridSearchCV(smote_pipe_rf, param_grid, cv=10,
                      scoring="average_precision")
search.fit(X_train, y_train)
results = pd.DataFrame(search.cv_results_)
results.plot("param_smote__k_neighbors", ["mean_test_score", "mean_train_score"])
```



# Summary

- Always check roc\_auc and AP, look at curves
- Undersampling is very fast and can help!
- Undersampling + Ensembles worth a try.
- Many smart sampling strategies, mixed outcomes
- SMOTE allows adding new interpolated samples
- Mixed outcomes with SMOTE, also definition a bit unclear

# References

- <https://arxiv.org/pdf/1106.1813.pdf>
- <http://cs.nju.edu.cn/zhouzh/zhouzh.files/publication/tsmcb09.pdf>

# Miscellaneous